

Topic 1: Introduction¹

(Version of 7th November 2025)

Pierre Flener and Jean-Noël Monette

Optimisation Group
Department of Information Technology
Uppsala University
Sweden

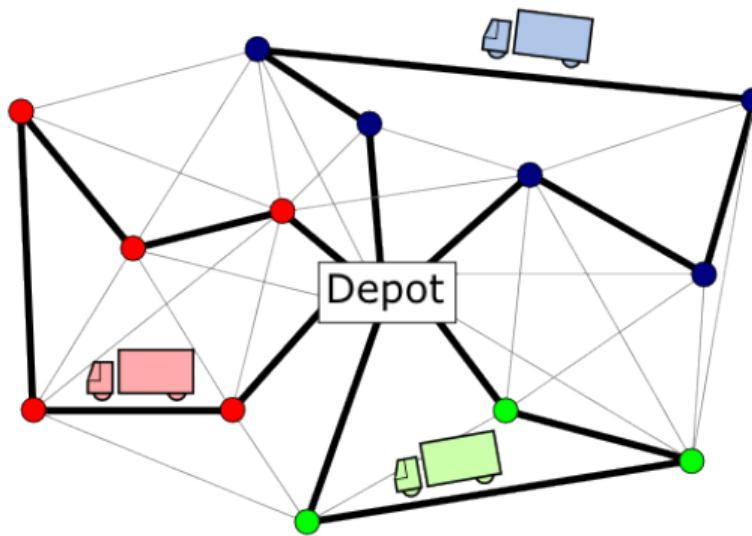
Course 1DL451:
Modelling for Combinatorial Optimisation

¹Based partly on material by Guido Tack



Optimisation

- Constraint Problems
- Combinatorial Optimisation
- Modelling (in MiniZinc)
- Solving
- The MiniZinc Toolchain
- Course Information



Optimisation is a science of **service**:
to scientists, to engineers, to artists, and to society.



Slogan of the Course

Constraint
Problems

Combinatorial
Optimisation

Modelling
(in MiniZinc)

Solving

The MiniZinc
Toolchain

Course
Information

How to solve problems
without knowing how to solve problems?



MiniZinc Challenge 2024: 15 (of 20) Problems and Winners

Problem and Model	Backend and Solver	Solving Technology
aircraft-disassembly	CP-SAT (by Google)	portfolio: LCG, MIP, CBLS
cable-tree-wiring	CP-SAT (by Google)	portfolio: LCG, MIP, CBLS
community-detection	MZN/Gurobi	MIP
compression	PicatSAT	SAT
concert-hall-cap	MZN/Gurobi	MIP
fox-geese-corn	MZN/Gurobi	MIP
hoist-benchmark	Chuffed	hybrid: LCG = CP + SAT
monitor-placement-1id	CP Optimizer (by IBM)	CP
neighbours	CP-SAT (by Google)	portfolio: LCG, MIP, CBLS
peaceable_queens	Gecode-Dexter	portfolio: CP, LNS
portal	Chuffed	hybrid: LCG = CP + SAT
tiny-cvrp	MZN/Gurobi	MIP
train-scheduling	CP-SAT (by Google)	portfolio: LCG, MIP, CBLS
word-equations	PicatSAT	SAT
yumi-dynamic	CP-SAT (by Google)	portfolio: LCG, MIP, CBLS



Outline

1. Constraint Problems

2. Combinatorial Optimisation

3. Modelling (in MiniZinc)

4. Solving

5. The MiniZinc Toolchain

6. Course Information



Outline

1. Constraint Problems

2. Combinatorial Optimisation

3. Modelling (in MiniZinc)

4. Solving

5. The MiniZinc Toolchain

6. Course Information



Example (Agricultural experiment design)

	plot1	plot2	plot3	plot4	plot5	plot6	plot7
barley							
corn							
millet							
oats							
rye							
spelt							
wheat							

Constraints to be satisfied:

- 1 Equal growth load: Every plot grows 3 grains.
- 2 Equal sample size: Every grain is grown in 3 plots.
- 3 Balance: Every grain pair is grown in 1 common plot.

Instance: 7 plots, 7 grains, 3 grains/plot, 3 plots/grain, balance 1.



Example (Agricultural experiment design)

	plot1	plot2	plot3	plot4	plot5	plot6	plot7
barley	✓	✓	✓	—	—	—	—
corn	✓	—	—	✓	✓	—	—
millet	✓	—	—	—	—	✓	✓
oats	—	✓	—	✓	—	✓	—
rye	—	✓	—	—	✓	—	✓
spelt	—	—	✓	✓	—	—	✓
wheat	—	—	✓	—	✓	✓	—

Constraints to be satisfied:

- 1 Equal growth load: Every plot grows 3 grains.
- 2 Equal sample size: Every grain is grown in 3 plots.
- 3 Balance: Every grain pair is grown in 1 common plot.

Instance: 7 plots, 7 grains, 3 grains/plot, 3 plots/grain, balance 1.



Example (Doctor rostering)

	Mon	Tue	Wed	Thu	Fri	Sat	Sun
Doctor A							
Doctor B							
Doctor C							
Doctor D							
Doctor E							

Constraints to be satisfied:

- 1 #on-call doctors / day = 1
- 2 #operating doctors / weekday ≤ 2
- 3 #operating doctors / week ≥ 7
- 4 #appointed doctors / week ≥ 4
- 5 day off after operation day
- 6 ...

Objective function to be minimised: Cost: ...



Example (Doctor rostering)

	Mon	Tue	Wed	Thu	Fri	Sat	Sun
Doctor A	call	none	oper	none	oper	none	none
Doctor B	appt	call	none	oper	none	none	call
Doctor C	oper	none	call	appt	appt	call	none
Doctor D	appt	oper	none	call	oper	none	none
Doctor E	oper	none	oper	none	call	none	none

Constraints to be satisfied:

- 1 #on-call doctors / day = 1
- 2 #operating doctors / weekday ≤ 2
- 3 #operating doctors / week ≥ 7
- 4 #appointed doctors / week ≥ 4
- 5 day off after operation day
- 6 ...



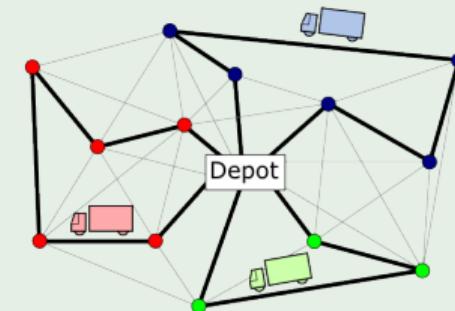
Objective function to be minimised: Cost: ...

Example (Vehicle routing: parcel delivery)

Given a depot with parcels for clients and a vehicle fleet,
find which vehicle visits which client when.

Constraints to be **satisfied**:

- 1 All parcels are delivered on time.
- 2 No vehicle is overloaded.
- 3 Driver regulations are respected.
- 4 ...



Objective function to be **minimised**:

- Cost: the total fuel consumption and driver salary.

Example (Travelling salesperson: optimisation TSP)

Given a map and cities,
find a **shortest** route visiting each city once and returning to the starting city.



Applications in Air Traffic Management

Constraint
Problems

Combinatorial
Optimisation

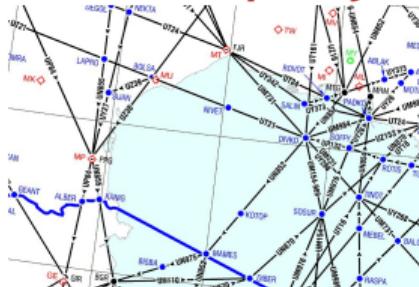
Modelling
(in MiniZinc)

Solving

The MiniZinc
Toolchain

Course
Information

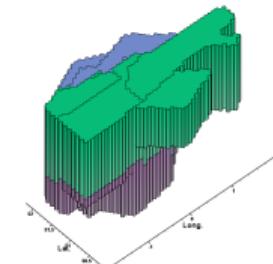
Demand vs capacity



Contingency planning

Flow	Time Span	Hourly Rate
From: Arlanda	00:00 – 09:00	3
	09:00 – 18:00	5
	18:00 – 24:00	2
From: Arlanda	00:00 – 12:00	4
	12:00 – 24:00	3
...

Airspace sectorisation



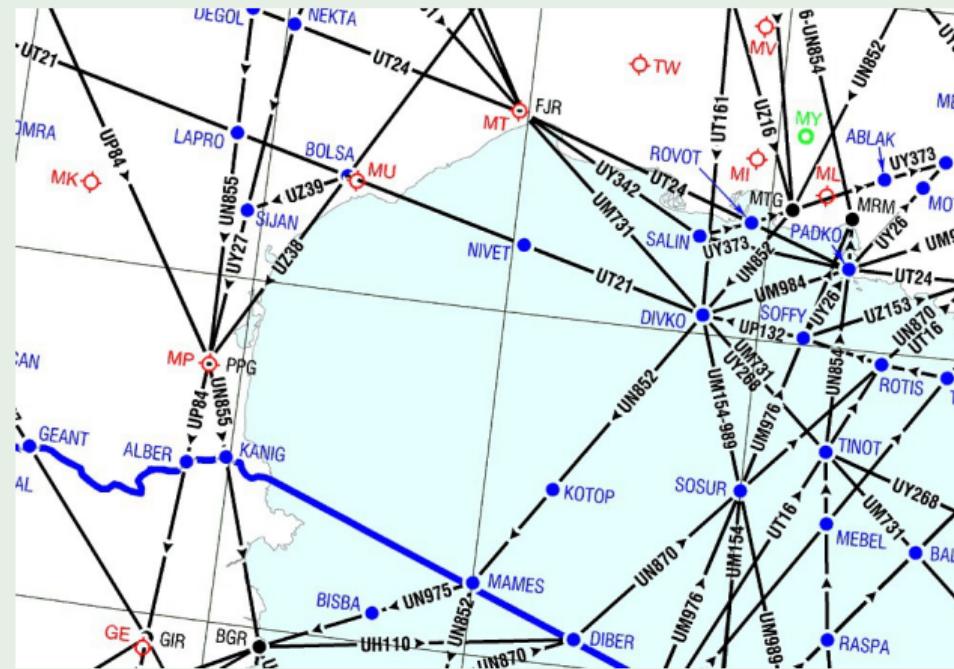
Workload balancing





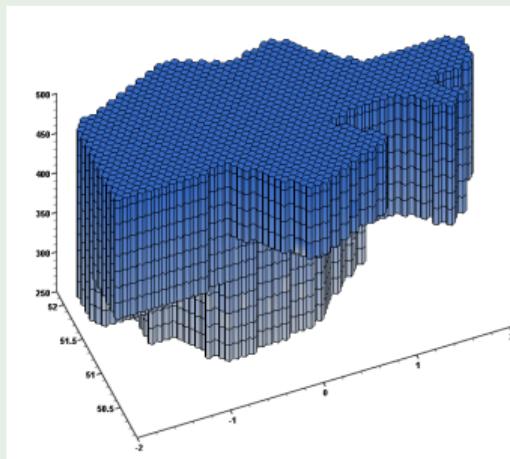
Example (Air-traffic demand-capacity balancing)

Reroute flights, in height and speed, so as to balance the workload of air traffic controllers in a multi-sector airspace:

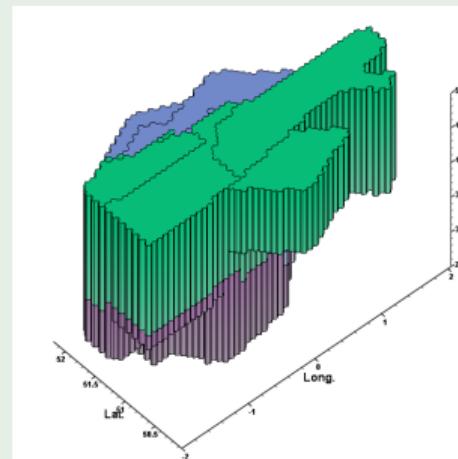


Example (Airspace sectorisation)

Given an airspace split into c cells, a targeted number s of sectors, and flight schedules.



Find a colouring of the c cells into s connected convex sectors, with minimal imbalance of the workloads of their air traffic controllers.



There are s^c possible colourings, but very few optimally satisfy the constraints: is **intelligent** search necessary?



Applications in Biology and Medicine

Constraint
Problems

Combinatorial
Optimisation

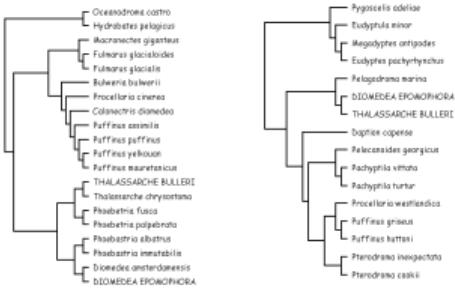
Modelling
(in MiniZinc)

Solving

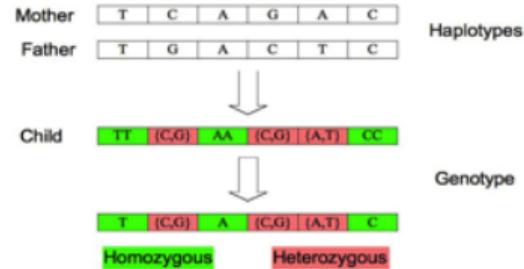
The MiniZinc
Toolchain

Course
Information

Phylogenetic supertree



Haplotype inference



Medical image analysis

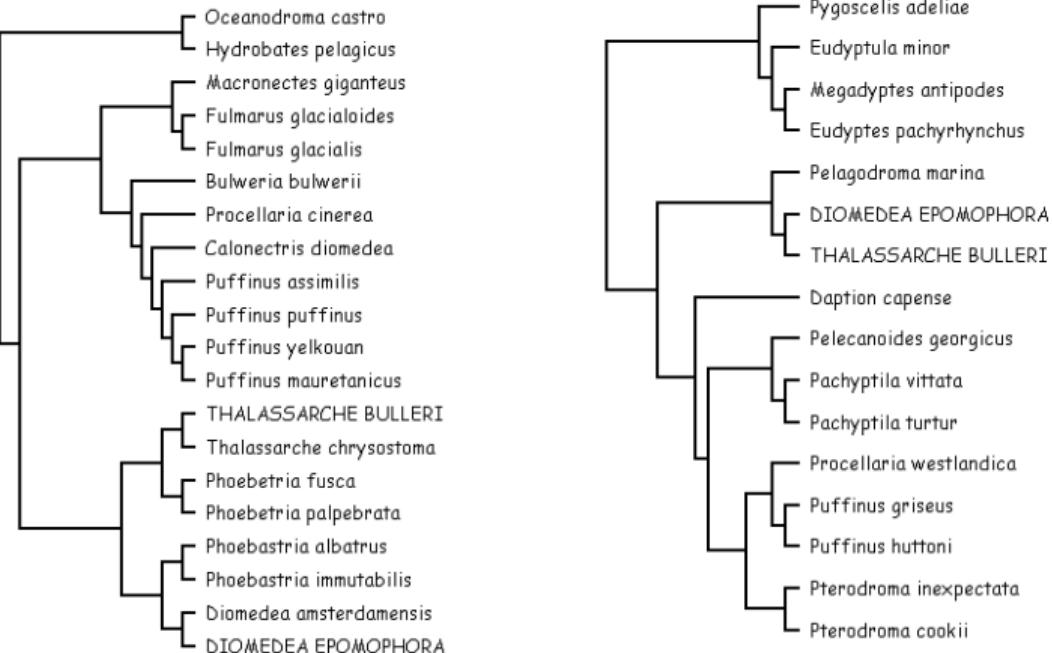


Doctor rostering





Example (What supertree is maximally consistent with several given trees that share some species?)





Example (Haplotype inference by pure parsimony)

Given n child genotypes, with homo- and heterozygous sites:

A	C / G	T	C	A / T	C
...					
A / T	G	T	C / G	A	C
...					

find a minimal set of (at most $2 \cdot n$) parent haplotypes:

A	C	T	C	T	C
...					
A	G	T	C	A	C
...					
T	G	T	G	A	C
...					

so that each given genotype conflates 2 found haplotypes.



Applications in Programming and Testing

Constraint
Problems

Combinatorial
Optimisation

Modelling
(in MiniZinc)

Solving

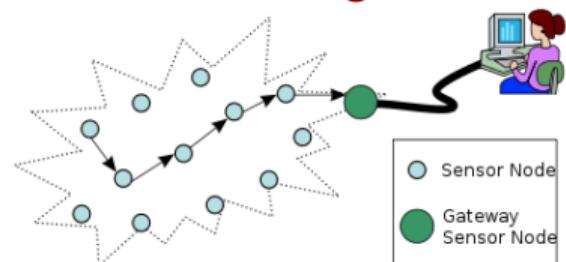
The MiniZinc
Toolchain

Course
Information

Robot programming



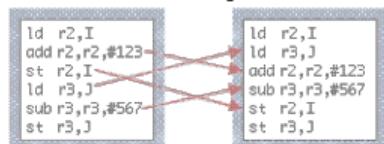
Sensor-net configuration



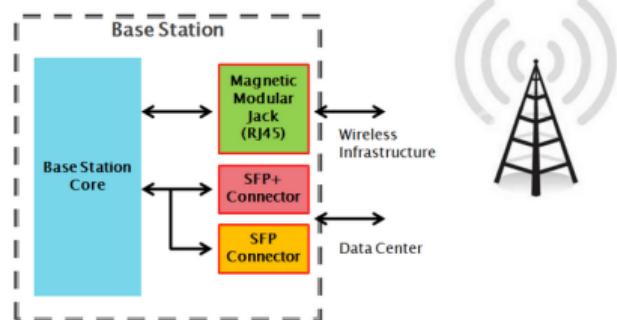
Compiler design

COMPILERS
FOR INSTRUCTION SCHEDULING

C Compiler
C++ Compiler



Base-station testing





Other Application Areas

Constraint
Problems

Combinatorial
Optimisation

Modelling
(in MiniZinc)

Solving

The MiniZinc
Toolchain

Course
Information

School timetabling

	Monday	Tuesday	Wednesday	Thursday	Friday
9:00	KMF200 Ordinary Differential Equations I FTII		LNUC201 Computer Graphics (I) Dual	KMF202 Numerical Analysis I Williamsson, GÖS	
10:00				KMF200 Ordinary Differential Equations I FTII	KMF200 Ordinary Differential Equations I KÖS
11:00	KMF200 Ordinary Differential Equations I KÖS, Room: 23		LNUC201 Computer Graphics (I) Dual	KMF202 Numerical Analysis I Williamsson, GÖS	KMF200 Ordinary Differential Equations I KÖS
12:00	CSD102 Algorithm and Data Structures I		KMF202 Further Linear Algebra IIS	KMF202 Ordinary Differential Equations I Williamsson, GÖS	
13:00		KMF200 Further Linear Algebra IIS	CSD102 Computer Graphics I		KMF202 Further Linear Algebra IIS
14:00				KMF202 Further Linear Algebra IIS	KMF202 Further Linear Algebra IIS
15:00			PASS Peer-Assisted Study IISY (LTF10 - LTF12 - IISY)		
16:00	CSD102 Computer Graphics I			KMF202 Further Linear Algebra IIS	
17:00		CSD102 Algorithm and Data Structures I			
18:00		CSD102 Algorithm and Data Structures I			

Security: SQL injection?



Sports tournament design

svensk handball



Container packing





Definitions

In a **constraint problem**, values have to be **found** for all the unknowns, called **variables** (in the mathematical sense; also called **decision variables**) and ranging over **given** sets, called **domains**, so that:

- All the given **constraints** on the decision variables are **satisfied**.
- Optionally: A given **objective function** on the decision variables has an optimal value: either a **minimal** cost or a **maximal** profit.

A **candidate solution** to a constraint problem maps each decision variable to a value within its domain; it is:

- **feasible** if all the constraints are satisfied;
- **optimal** if the objective function takes an optimal value.

The **search space** consists of all candidate solutions.

A **solution** to a **satisfaction problem** is feasible.

An **optimal solution** to an **optimisation problem** is feasible and optimal.



This is one of the seven **Millennium Prize** problems of the Clay Mathematics Institute (Massachusetts, USA), each worth 1 million US\$.

Informally:

- P = class of problems that need **no** search to be solved
NP = class of problems that **might** need search to solve
- P = class of problems with **easy-to-compute** solutions
NP = class of problems with **easy-to-check** solutions

Thus: Can search always be avoided ($P = NP$),
or is search sometimes necessary ($P \neq NP$)?

Problems that are solvable in polynomial time (in the input size) are considered **tractable**, aka **easy**.

Problems needing super-polynomial time are considered **intractable**, aka **hard**.



NP Completeness: Examples

Given a digraph (V, E) :

Examples

- Finding a **shortest path** takes $\mathcal{O}(V \cdot E)$ time and is thus in P.
- Determining the existence of a simple path (which has distinct vertices), from a given single source, that has *at least* a given number ℓ of edges is NP-complete. Hence finding a **longest path** seems hard:
increase ℓ starting from a trivial lower bound, until answer is ‘no’.

Examples

- Finding an **Euler tour** (which visits each *edge* once) takes $\mathcal{O}(E)$ time and is thus in P.
- Determining the existence of a **Hamiltonian cycle** (which visits each *vertex* once) is NP-complete.



NP Completeness: More Examples

Examples

- ***n-SAT***: Determining the satisfiability of a conjunction of disjunctions of n Boolean literals is in P for $n = 2$ but NP-complete for $n = 3$.
- ***SAT***: Determining the satisfiability of a formula over Boolean literals is NP-complete.
- ***Clique***: Determining the existence of a clique (complete subgraph) of a given size in a graph is NP-complete.
- ***Vertex Cover***: Determining the existence of a vertex cover (a vertex subset with at least one endpoint for all edges) of a given size in a graph is NP-complete.
- ***Subset Sum***: Determining the existence of a subset, of a given set, that has a given sum is NP-complete.



Search spaces are often larger than the universe!



Many important real-life problems are NP-hard or worse: their real-life instances can only be solved exactly and fast enough by **intelligent** search, unless $P = NP$. **NP-hardness is not where the fun ends, but where it begins!**



Example (Optimisation TSP over n cities)

A brute-force algorithm evaluates all $n!$ candidate routes:

- A computer of today evaluates 10^6 routes / second:

n	time
11	40 seconds
14	1 day
18	203 years
20	77k years

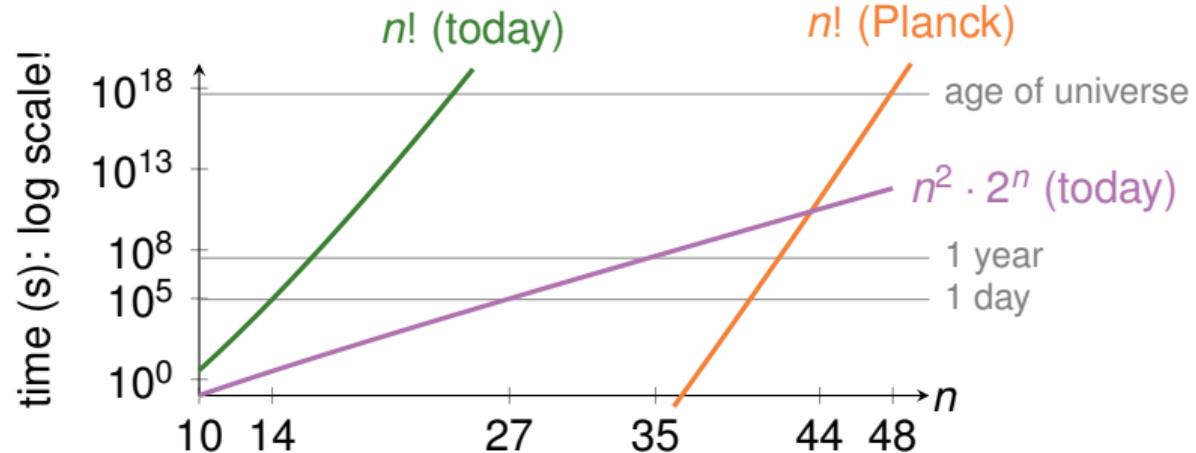
- Planck time is shortest useful interval: $\approx 5.4 \cdot 10^{-44}$ second; a Planck computer would evaluate $1.8 \cdot 10^{43}$ routes / second:

n	time
37	0.7 seconds
41	20 days
48	1.5 · age of universe

The dynamic program by Bellman-Held-Karp “only” takes $\mathcal{O}(n^2 \cdot 2^n)$ time: a computer of today takes a day for $n = 27$, a year for $n = 35$, the age of the universe for $n = 67$, and beats the $\mathcal{O}(n!)$ algo on Planck computer for $n \geq 44$.

Intelligent Search upon NP-Hardness

Do not give up but try to stay ahead of the curve:
there is an instance size until which an **exact** algorithm is fast enough!



Concorde TSP Solver beats the Bellman-Held-Karp exact algo: it uses local search & approximation algos, but sometimes proves exactness of its optima.
The largest instance solved exactly, in 136 CPU years in 2006, has $n = 85900$.



Outline

1. Constraint Problems

2. Combinatorial Optimisation

3. Modelling (in MiniZinc)

4. Solving

5. The MiniZinc Toolchain

6. Course Information



A **solving technology** offers languages, methods, and tools for:

what: **Modelling** constraint problems in a **declarative** language.

and / or

how: **Solving** constraint problems **intelligently**:

- **Search**: Explore the space of candidate solutions.
- **Reasoning**: Reduce the space of candidate solutions.
- **Relaxation**: Exploit solutions to easier problems.

A **solver** is a program that takes a model and data as input and tries to solve that problem instance.

Combinatorial (= discrete) optimisation covers satisfaction *and* optimisation problems for variables ranging over *discrete* sets: **combinatorial problems**.

The ideas in this course extend to continuous optimisation, to soft optimisation, and to stochastic optimisation.



Examples (Solving technologies)

With general-purpose solvers, taking model and data as input:

- Boolean satisfiability (SAT)
- SAT (resp. optimisation) modulo theories (SMT and OMT)
- (Mixed) integer linear programming (IP and MIP)
- Constraint programming (CP) ☞ available via [1DL705](#)
- ...
- Hybrid technologies ($LCG = CP + SAT, \dots$) and portfolios

Methodologies, *usually without* modelling and solvers:

- Dynamic programming (DP)
- Greedy algorithms
- Approximation algorithms
- Local search (LS)
- ...



Outline

1. Constraint Problems

2. Combinatorial Optimisation

3. Modelling (in MiniZinc)

4. Solving

5. The MiniZinc Toolchain

6. Course Information



What vs How

Example

Consider the **problem** of sorting an array A of n numbers into an array S of increasing-or-equal numbers.

A **formal specification** is:

$$\text{sort}(A, S) \equiv \text{permutation}(A, S) \wedge \text{increasing}(S)$$

saying that S must be a permutation of A in increasing order.

Seen as a generate-and-test **algorithm**, it takes $\mathcal{O}(n!)$ time, but it can be refined into the existing $\mathcal{O}(n \log n)$ algorithms.

A **specification** is a **declarative** description of **what** problem is to be solved. An **algorithm** is a **procedural** description of **how** to solve the problem (fast).



Modelling vs Programming

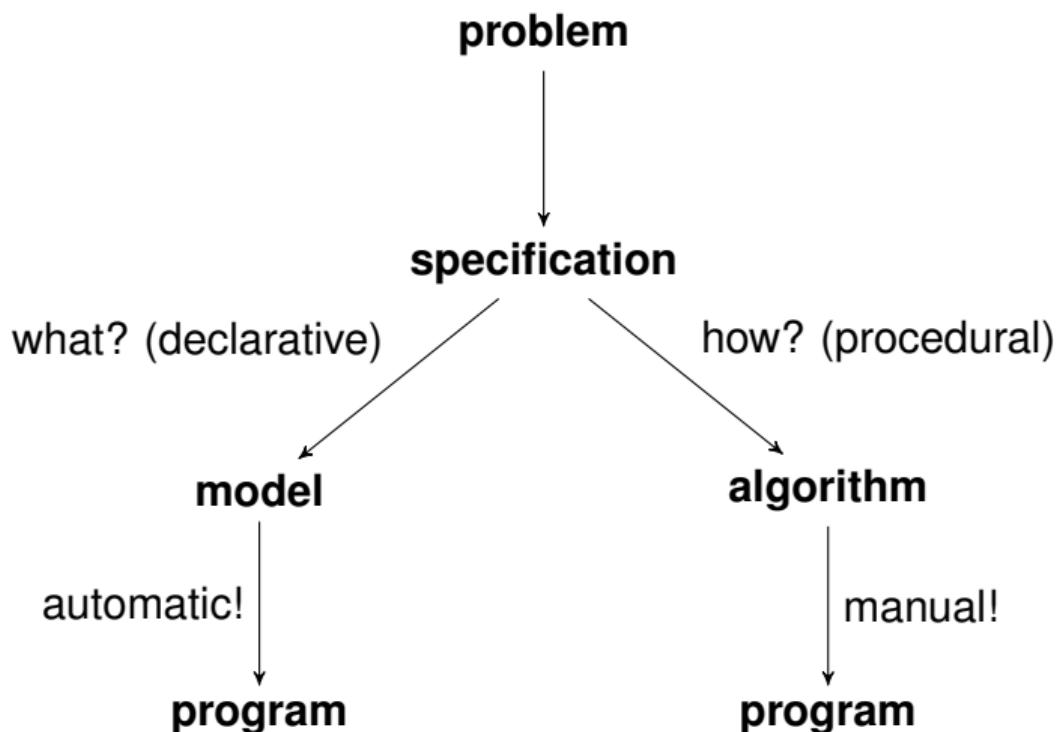
Constraint
Problems

Combinatorial
Optimisation

Modelling
(in MiniZinc)

Solving
The MiniZinc
Toolchain

Course
Information





Example (Sudoku)

Constraint
ProblemsCombinatorial
OptimisationModelling
(in MiniZinc)

Solving

The MiniZinc
ToolchainCourse
Information

8								
	3	6						
7		9		2				
5			7					
		4	5	7				
		1			3			
	1				6	8		
	8	5			1			
9				4				

8	1	2	7	5	3	6	4	9
9	4	3	6	8	2	1	7	5
6	7	5	4	9	1	2	8	3
1	5	4	2	3	7	8	9	6
3	6	9	8	4	5	7	2	1
2	8	7	1	6	9	5	3	4
5	2	1	9	7	4	3	6	8
4	3	8	5	2	6	9	1	7
7	9	6	3	1	8	4	5	2

A Sudoku is a 9-by-9 array of integers in the range 1..9. Some of the elements are provided as parameters. The remaining elements are unknowns that have to satisfy the following constraints:

- 1 the elements in each row are all different;
- 2 the elements in each column are all different;
- 3 the elements in each 3-by-3 block are all different.



Example (Sudoku)

Google

Translate

Turn off instant

English Turkish Swedish English - detected ▾



MiniZinc Turkish Swedish ▾

Translate

A Sudoku is a 9-by-9 array of integers in the interval 1..9. ×
Some of the elements are provided as parameters.

The remaining elements are unknowns
that have to satisfy the following constraints:
- the elements in each row are all different;
- the elements in each column are all different;
- the elements in each 3-by-3 block are all different.

```
array[1..9,1..9] of var 1..9: Sudoku;
solve satisfy;
constraint forall(row in 1..9)
  (alldifferent(Sudoku[row, ..]));
constraint forall(col in 1..9)
  (alldifferent(Sudoku[.., col]));
constraint forall(i,j in {0,3,6})
  (alldifferent(Sudoku[i+1..i+3, j+1..j+3]));
```



Example (Sudoku ↗)

8								
	3	6						
7			9		2			
5				7				
			4	5	7			
		1				3		
	1					6	8	
	8	5				1		
9				4				

8	1	2	7	5	3	6	4	9
9	4	3	6	8	2	1	7	5
6	7	5	4	9	1	2	8	3
1	5	4	2	3	7	8	9	6
3	6	9	8	4	5	7	2	1
2	8	7	1	6	9	5	3	4
5	2	1	9	7	4	3	6	8
4	3	8	5	2	6	9	1	7
7	9	6	3	1	8	4	5	2

```
-2 array[1..9,1..9] of var 1..9: Sudoku;
-1 ... % load the hints
0 solve satisfy;
1 constraint forall(row in 1..9) (all_different(Sudoku[row,...]));
2 constraint forall(col in 1..9) (all_different(Sudoku[...,col]));
3 constraint forall(i,j in {0,3,6})
    (all_different(Sudoku[i+1..i+3,j+1..j+3]));
```



Example (Agricultural experiment design, AED)

	plot1	plot2	plot3	plot4	plot5	plot6	plot7
barley	✓	✓	✓	—	—	—	—
corn	✓	—	—	✓	✓	—	—
millet	✓	—	—	—	—	✓	✓
oats	—	✓	—	✓	—	✓	—
rye	—	✓	—	—	✓	—	✓
spelt	—	—	✓	✓	—	—	✓
wheat	—	—	✓	—	✓	✓	—

Constraints to be satisfied:

- 1 Equal growth load: Every plot grows 3 grains.
- 2 Equal sample size: Every grain is grown in 3 plots.
- 3 Balance: Every grain pair is grown in 1 common plot.

Instance: 7 plots, 7 grains, 3 grains/plot, 3 plots/grain, balance 1.

General term: balanced incomplete block design (BIBD).



Example (Agricultural experiment design, AED)

	plot1	plot2	plot3	plot4	plot5	plot6	plot7
barley	1	1	1	0	0	0	0
corn	1	0	0	1	1	0	0
millet	1	0	0	0	0	1	1
oats	0	1	0	1	0	1	0
rye	0	1	0	0	1	0	1
spelt	0	0	1	1	0	0	1
wheat	0	0	1	0	1	1	0

Constraints to be satisfied:

- 1 Equal growth load: Every plot grows 3 grains.
- 2 Equal sample size: Every grain is grown in 3 plots.
- 3 Balance: Every grain pair is grown in 1 common plot.

Instance: 7 plots, 7 grains, 3 grains/plot, 3 plots/grain, balance 1.

General term: balanced incomplete block design (BIBD).



In a BIBD, the plots are called **blocks** and the grains are called **varieties**:

Example (BIBD *integer* model ↗: ✓ ↵ 1 and – ↵ 0)

```
-3 enum Varieties; enum Blocks;
-2 int: blockSize; int: sampleSize; int: balance;
-1 array[Varieties,Blocks] of var 0..1: BIBD; % BIBD[v,b]=1 iff v is in b
0 solve satisfy;
1 constraint forall(b in Blocks) (blockSize = sum(BIBD[.,b]));
2 constraint forall(v in Varieties) (sampleSize = sum(BIBD[v,.]));
3 constraint forall(v, w in Varieties where v < w)
    (balance = sum([BIBD[v,b]*BIBD[w,b] | b in Blocks]));
```

Example (Instance data for our AED ↗)

```
-3 Varieties = {barley,...,wheat}; Blocks = {plot1,...,plot7};
-2 blockSize = 3; sampleSize = 3; balance = 1;
```

Using the `count` abstraction instead of `sum`:Example (BIBD *integer* model ↗: ✓ ↵ 1 and – ↵ 0)

```
-3 enum Varieties; enum Blocks;
-2 int: blockSize; int: sampleSize; int: balance;
-1 array[Varieties,Blocks] of var 0..1: BIBD; % BIBD[v,b]=1 iff v is in b
0 solve satisfy;
1 constraint forall(b in Blocks) (blockSize = count(BIBD[.,b], 1));
2 constraint forall(v in Varieties) (sampleSize = count(BIBD[v,.], 1));
3 constraint forall(v, w in Varieties where v < w)
    (balance = count([BIBD[v,b]*BIBD[w,b] | b in Blocks], 1));
```

Example (Instance data for our AED ↗)

```
-3 Varieties = {barley,...,wheat}; Blocks = {plot1,...,plot7};
-2 blockSize = 3; sampleSize = 3; balance = 1;
```

Using the `count` abstraction over linear expressions:Example (BIBD *integer* model ↗: ✓ ↵ 1 and – ↵ 0)

```
-3 enum Varieties; enum Blocks;
-2 int: blockSize; int: sampleSize; int: balance;
-1 array[Varieties,Blocks] of var 0..1: BIBD; % BIBD[v,b]=1 iff v is in b
0 solve satisfy;
1 constraint forall(b in Blocks) (blockSize = count(BIBD[.,b], 1));
2 constraint forall(v in Varieties) (sampleSize = count(BIBD[v,.], 1));
3 constraint forall(v, w in Varieties where v < w)
    (balance = count([BIBD[v,b]+BIBD[w,b] | b in Blocks], 2));
```

Example (Instance data for our AED ↗)

```
-3 Varieties = {barley,...,wheat}; Blocks = {plot1,...,plot7};
-2 blockSize = 3; sampleSize = 3; balance = 1;
```



Reconsider the model fragment:

```
2 constraint forall(v in Varieties) (sampleSize = count(BIBD[v,..], 1));
```

This constraint is **declarative** (and by the way non-linear),
so read it using only the verb “to be” or synonyms thereof:

*for all varieties v,
the count of occurrences of 1 in row v of BIBD
must equal sampleSize*

The constraint is **not procedural**:

*for all varieties v,
we first count the occurrences of 1 in row v
and then check if that count equals sampleSize*

The latter reading is appropriate for solution **checking**,
but solution **finding** usually performs no such procedural counting.



Example (Idea for another BIBD model)

barley	{plot1, plot2, plot3}	}
corn	{plot1, plot4, plot5}	}
millet	{plot1, plot6, plot7}	}
oats	{plot2, plot4, plot6}	}
rye	{plot2, plot5, plot7}	}
spelt	{plot3, plot4, plot7}	}
wheat	{plot3, plot5, plot6}	}

Constraints to be satisfied:

- 1 Equal growth load: Every plot grows 3 grains.
- 2 Equal sample size: Every grain is grown in 3 plots.
- 3 Balance: Every grain pair is grown in 1 common plot.



Example (BIBD set model ↗: a block set per variety)

```
-3 enum Varieties; enum Blocks;
-2 int: blockSize; int: sampleSize; int: balance;
-1 array[Varieties] of var set of Blocks: BIBD; % BIBD[v] = blocks for v
0 solve satisfy;
1 constraint forall(b in Blocks)
    (blockSize = sum(v in Varieties) (b in BIBD[v]));
2 constraint forall(v in Varieties)
    (sampleSize = card(BIBD[v]));
3 constraint forall(v, w in Varieties where v < w)
    (balance      = card(BIBD[v] intersect BIBD[w]));
```

Example (Instance data for our AED ↗)

```
-3 Varieties = {barley, ..., wheat}; Blocks = {plot1, ..., plot7};
-2 blockSize = 3; sampleSize = 3; balance = 1;
```



Example (Doctor rostering)

	Mon	Tue	Wed	Thu	Fri	Sat	Sun
Doctor A	call	none	oper	none	oper	none	none
Doctor B	appt	call	none	oper	none	none	call
Doctor C	oper	none	call	appt	appt	call	none
Doctor D	appt	oper	none	call	oper	none	none
Doctor E	oper	none	oper	none	call	none	none

Constraints to be satisfied:

- 1 #on-call doctors / day = 1
- 2 #operating doctors / weekday ≤ 2
- 3 #operating doctors / week ≥ 7
- 4 #appointed doctors / week ≥ 4
- 5 day off after operation day
- 6 ...



Objective function to be minimised: Cost: ...



Example (Doctor rostering ↗)

```
-5 set of int: Days; % d mod 7 = 1 iff d is a Monday
-4 enum Doctors;
-3 enum ShiftTypes = {appt, call, oper, none};
-2 % Roster[i,j] = shift type of Dr i on day j:
-1 array[Doctors,Days] of var ShiftTypes: Roster;
0 solve minimize ...; % plug in an objective function
1 constraint forall(d in Days) (count(Roster[..,d],call) = 1);
2 constraint forall(d in Days where d mod 7 in 1..5)
    (count(Roster[..,d],oper) <= 2);
3 constraint count(Roster,oper) >= 7;
4 constraint count(Roster,appt) >= 4;
5 constraint forall(d in Doctors)
    (regular(Roster[d,..],"((oper none) | appt | call | none)*"));
6 ... % other constraints
```

Example (Instance data for our small hospital unit ↗)

```
-5 Days = 1..7;
-4 Doctors = {Dr_A, Dr_B, Dr_C, Dr_D, Dr_E};
```



Using decision variables as indices within arrays: **black magic?**

Example (Job allocation at minimal salary cost)

Given jobs `Jobs` and the salaries of work applicants `Apps`,
find a work applicant for each job

such that some constraints (on the qualifications of the work applicants for the jobs, on workload distribution, etc) are satisfied and the total salary cost is minimal:

```
1 array[Apps] of 0..1000: Salary; % Salary[a] = cost per job to appl. a
2 array[Jobs] of var Apps: Worker; % Worker[j] = appl. allocated job j
3 solve minimize sum(j in Jobs)(Salary[Worker[j]]); 
4 constraint ...; % qualifications, workload, etc
```



Using decision variables as indices within arrays: black magic?

Example (Vehicle routing: backbone model)

```
enum Cities = {AMS, BRU, LUX, CDG}
```

Next:

AMS	BRU	LUX	CDG





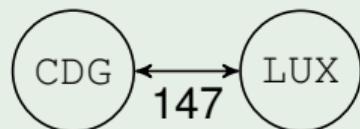
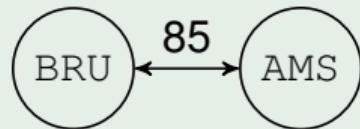
Using decision variables as indices within arrays: black magic?

Example (Vehicle routing: backbone model)

```
enum Cities = {AMS, BRU, LUX, CDG}
```

	AMS	BRU	LUX	CDG
Next:	BRU	AMS	CDG	LUX

So `all_different(Next)` is too weak!





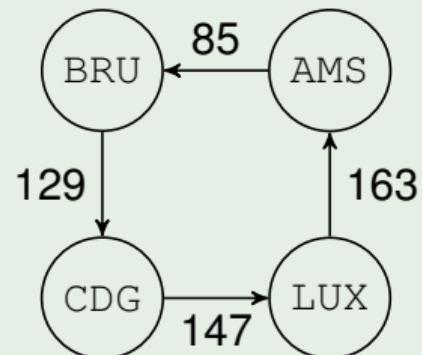
Using decision variables as indices within arrays: black magic?

Example (Vehicle routing: backbone model)

```
enum Cities = {AMS, BRU, LUX, CDG}
```

	AMS	BRU	LUX	CDG
Next:	BRU	CDG	AMS	LUX

Let us use `circuit` (`Next`) instead:





Using decision variables as indices within arrays: black magic?

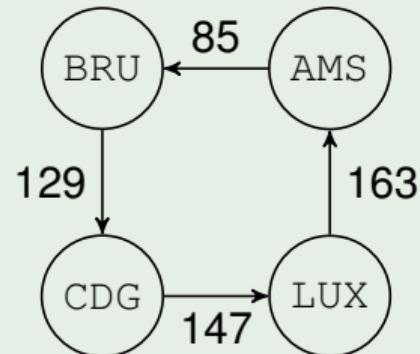
Example (Vehicle routing: backbone model)

```
enum Cities = {AMS, BRU, LUX, CDG}
```

	AMS	BRU	LUX	CDG
Next:	BRU	CDG	AMS	LUX

Let us use `circuit(Next)` instead:

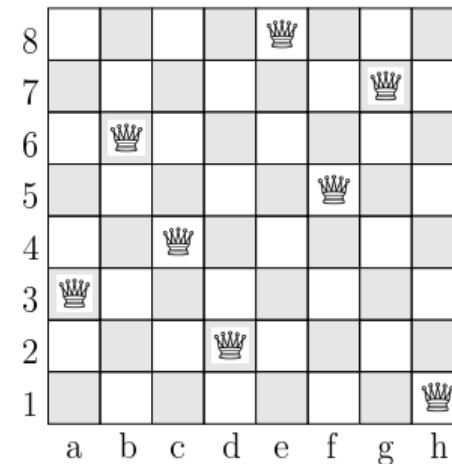
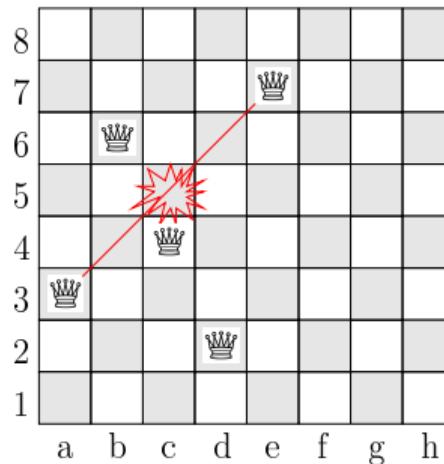
```
1 array[Cities,Cities] of int: Dist; % Dist[i,j] = distance from i to j
2 array[Cities] of var Cities: Next; % travel from c to Next[c]
3 solve minimize sum(c in Cities) (Dist[c,Next[c]]);
4 constraint circuit(Next);
5 constraint ...; % side constraints, if any
```





Toy Example: 8-Queens

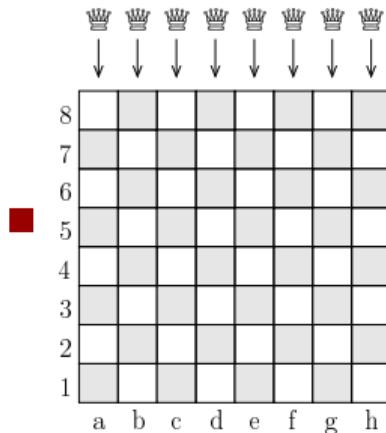
Can one place 8 queens onto an 8×8 chessboard so that all queens are in distinct rows, columns, and diagonals?





An 8-Queens Model

One of the many models, with **one** decision variable per queen:



Let **decision variable** $\text{Row}[c]$, of **domain** $1 \dots 8$, denote the row of the queen in column c , for c in $\{a, b, c, \dots, h\}$, which we rename into $1 \dots 8$. Example: $\text{Row}[3] = 4$ means that the queen of column 3 (column c in the picture) is in row 4. The **constraint** that all queens must be in distinct columns is **satisfied** by the choice of variables!

■ The remaining **constraints** to be **satisfied** are:

- All queens are in distinct rows: the var.s $\text{Row}[c]$ take distinct values for all c
- All queens are in distinct diagonals:
the expressions $\text{Row}[c] + c$ take distinct values for all c
the expressions $\text{Row}[c] - c$ take distinct values for all c



An 8-Queens Model in MiniZinc

Consider the following model [↗](#) in a file 8-queens.mzn:

```
1 include "globals.mzn"; % ensures that lines 4 to 6 are understood
2 int: n = 8; % the given number of queens
3 array[1..n] of var 1..n: Row; % Row[c] = the unknown row of the queen
      in column c;      enforces that all queens are in distinct columns
4 constraint all_different( Row ) ; % distinct rows
5 constraint all_different([Row[c]+c | c in 1..n]); % distinct up-dia.
6 constraint all_different([Row[c]-c | c in 1..n]); % distinct down-dia.
7 solve satisfy; % solve to satisfaction of all the constraints
8 output [show(Row)]; % pretty-printing of solutions
```

The `all_different(X)` constraint holds if and only if all the expressions in the array `X` take different values.



Modelling Concepts

- A **variable**, also called a **decision variable**, is an existentially quantified unknown of a problem.
- The **domain** of a decision variable x , here denoted by $\text{dom}(x)$, is the set of values in which x must take its value, if any.
- A **variable expression** takes a value that depends on the value of one or more decision variables.
- A **parameter** has a value from a problem description.
- Decision variables, parameters, and expressions are **typed**.

MiniZinc types are (arrays and sets of) Booleans, integers, floating-point numbers, enumerations, records, tuples, and strings, but not all these types can serve as types for decision variables.



Decision Variables, Parameters, and Identifiers

- Decision variables and parameters in a model are concepts very different from programming variables in an imperative or object-oriented program.
- A decision variable in a model is like a variable in mathematics: it is *not* given a value in a model or a formula, and its value is only fixed in a solution, if a solution exists.
- A parameter in a model must be given a value, but only once: we say that it is **instantiated**.
- A decision variable or parameter is referred to by an **identifier**.
- An **index identifier** of an array **comprehension** takes on all its designated values in turn. Example: the index `c` in the 8-queens model.



Parametric Models

- A parameter need not be instantiated inside a model.
Example: drop “=8” from “`int: n=8`” in the 8-queens model to make it an `n`-queens model, and rename `8-queens.mzn` into `n-queens.mzn`.
- Data are values for parameters given outside a model:
either in a datafile (.dzn suffix), or at the command line,
or interactively in the integrated development environment (IDE).
- A parametric model has uninstantiated parameters.
- An instance is a pair of a parametric model and data.



Modelling Concepts (end)

- A **constraint** is a restriction on the values that its decision variables can take together; equivalently, it is a Boolean-valued variable expression that must be true.
- An **objective function** is a numeric variable expression whose value is to be either minimised or maximised.
- An **objective** states what is being asked for:
 - find a first solution
 - find a solution minimising an objective function
 - find a solution maximising an objective function
 - find all solutions
 - count the number of solutions
 - prove that there is no solution
 - ...



Constraint-Based Modelling

MiniZinc is a high-level **constraint-based** modelling language (*not* a solver):

- There are several **types** for decision variables: `bool`, `int`, `float`, `enum`,
`string`, `tuple`, `record`, and `set`,
possibly as elements of multidimensional matrices (`array`).
- There is a large vocabulary of **predicates** (`<`, `<=`, `=`, `!=`, `>=`, `>`,
`all_different`, `circuit`, `regular`, ...), **functions** (`+`, `-`, `*`, `card`,
`count`, `intersect`, `sum`, ...), and **logical connectives & quantifiers**
(`not`, `/\`, `\/`, `->`, `<-`, `<->`, `forall`, `exists`, ...).
- There is support for *both* constraint **satisfaction** (`satisfy`)
and constrained **optimisation** (`minimize` and `maximize`).

Most modelling languages are (much) lower-level than this!



Correctness Is Not Enough for Models

Constraint
Problems

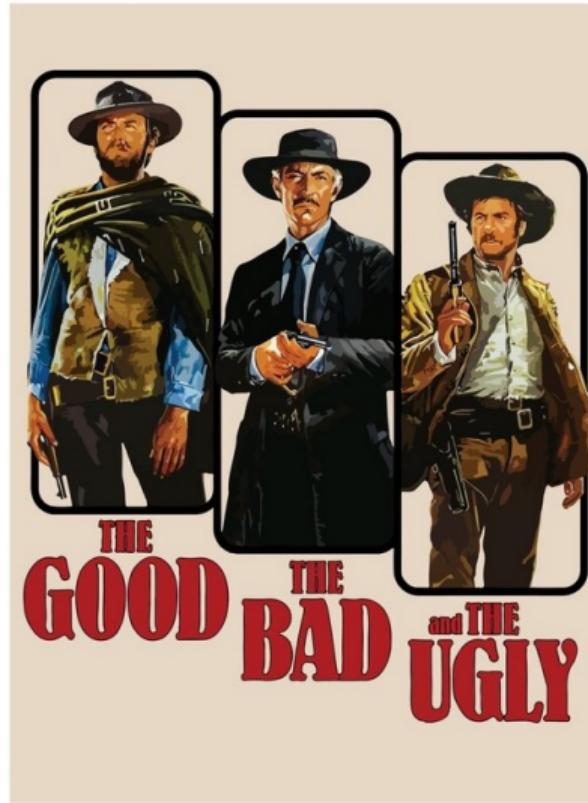
Combinatorial
Optimisation

Modelling
(in MiniZinc)

Solving

The MiniZinc
Toolchain

Course
Information





Modelling is an Art!

There are good and bad models for each constraint problem:

- Different models of a problem may take different time on the same solver for the same instance.
- Different models of a problem may scale differently on the same solver for instances of growing size.
- Different solvers may take different time on the same model for the same instance.

Good modellers are worth their weight in gold!

Use solvers: based on decades of cutting-edge research, they are very hard to beat on exact solving.



Outline

1. Constraint Problems

2. Combinatorial Optimisation

3. Modelling (in MiniZinc)

4. Solving

5. The MiniZinc Toolchain

6. Course Information



Solutions to a problem instance can be found by running a MiniZinc *backend*, that is a MiniZinc wrapper for a particular solver, on a file containing a model of the problem.

Example (Solving the 8-queens instance)

Let us run the solver Gecode, of CP technology, from the command line:

```
minizinc --solver gecode 8-queens.mzn
```

The result is printed on stdout:

```
[4, 2, 7, 3, 6, 8, 5, 1]
```

This means that the queen of column 1 is in row 4 (note that MiniZinc uses 1-based indexing), the queen of column 2 is in row 2, and so on.

Use the command-line flag `-a` to ask for all solutions:

the line ----- is printed after each solution, but the

line ===== is printed after the last (the 92nd here) solution.



Definition (Solving = Search + Reasoning + Relaxation)

- **Search**: Explore the space of candidate solutions.
- **Reasoning**: Reduce the space of candidate solutions.
- **Relaxation**: Exploit solutions to easier problems.

Definition (Systematic Search: guarantees ultimately exact solving)

Progressively build a solution, and backtrack if necessary.

Use **reasoning** and **relaxation** in order to reduce the **search** effort.

It is used in most SAT, SMT, OMT, CP, LCG, and MIP solvers.

Definition (Local Search: trades guarantee of exact solving for speed)

Start from a candidate solution and iteratively modify it a bit, until time-out.

It is the basic idea behind LS and genetic algorithm (GA) technologies.

For some details, see Topic 7: Solving Technologies.



There Are So Many Solving Technologies

- No technology universally dominates all the others.
- One should test several technologies on each problem.
- Some technologies have **no** modelling languages:
LS, DP, and GA are rather methodologies.
- Some technologies have **standardised** modelling languages across all solvers: SAT, SMT, OMT, and (M)IP.
- Some technologies have **non-standardised** modelling languages across their solvers: CP and LCG.



Model and Solve

Advantages:

- + Declarative model of a problem.
- + Easy adaptation to changing problem requirements.
- + Use of powerful solving technologies that are based on decades of cutting-edge research.

Disadvantages:

- Do I need to learn several modelling languages? **No!**
- Do I need to understand the used solving technologies in order to get the most out of them? **Yes, but . . . !**



Outline

1. Constraint Problems

2. Combinatorial Optimisation

3. Modelling (in MiniZinc)

4. Solving

5. The MiniZinc Toolchain

6. Course Information



MiniZinc

MiniZinc is a declarative language (*not a solver*)
for the constraint-based modelling of constraint problems:



- At Monash University, Australia
- Introduced in 2007; version 2.0 in 2014; version 3.0 imminent?
- Homepage: <https://www.minizinc.org>
- Integrated development environment (IDE)
- Annual **MiniZinc Challenge** for solvers, since 2008
- There are also **courses at Coursera**, also in Chinese



MiniZinc Features

- Declarative language for modelling **what** the problem is
- Separation of problem **model** and instance **data**
- **Open-source** toolchain
- Much **higher**-level language than those of (M)IP and SAT
- Solver-**independent** language
- Solving-technology-**independent** language
- Vocabulary of **predefined** types, predicates and functions
- Support for **user-defined** predicates and functions
- Support for annotations with hints on **how** to solve
- Ever-growing number of users, solvers, and other tools



MiniZinc Backends and Their Solvers

- SAT = Boolean satisfiability: Plingeling via PicatSAT, ...
- MIP = mixed integer programming: Cbc, FICO Xpress, Gurobi Optimizer, HiGHS, IBM ILOG CPLEX Optimizer, ...
- CP = constraint programming:
Choco, Gecode, JaCoP, Mistral, SICStus Prolog, ...
- CBLS = constraint-based LS (local search), without exactness guarantee:
Atlantis, OscaR.cbls via fzn-oscar-cbls, Yuck, ... : almost always time out
- LCG = lazy clause generation, a hybrid of CP and SAT: Chuffed, ...
- Other hybrid technologies: iZplus, MiniSAT(ID), SCIP, ...
- Portfolios: Google's CP-SAT of OR-Tools (with LCG, MIP, and LS), ...
- ..., SMT, OMT, ...



MiniZinc Backends and Their Solvers

- SAT = Boolean satisfiability: **Plingeling** via **PicatSAT**, ...
- MIP = mixed integer programming: **Cbc**, FICO Xpress, **Gurobi Optimizer**, **HiGHS**, IBM ILOG CPLEX Optimizer, ...
- CP = constraint programming:
Choco, **Gecode**, JaCoP, Mistral, SICStus Prolog, ...
- CBLS = constraint-based LS (local search), without exactness guarantee:
Atlantis, OscaR.cbls via fzn-oscar-cbils, **Yuck**, ... : almost always time out
- LCG = lazy clause generation, a hybrid of CP and SAT: **Chuffed**, ...
- Other hybrid technologies: iZplus, MiniSAT(ID), SCIP, ...
- Portfolios: Google's **CP-SAT** of OR-Tools (with LCG, MIP, and LS), ...
- ..., SMT, OMT, ...

The backends installed on the IT department's ThinLinc hardware are in red.
The commercial **Gurobi Optimizer** is under a free academic license:
you may **not** use it for non-academic purposes.

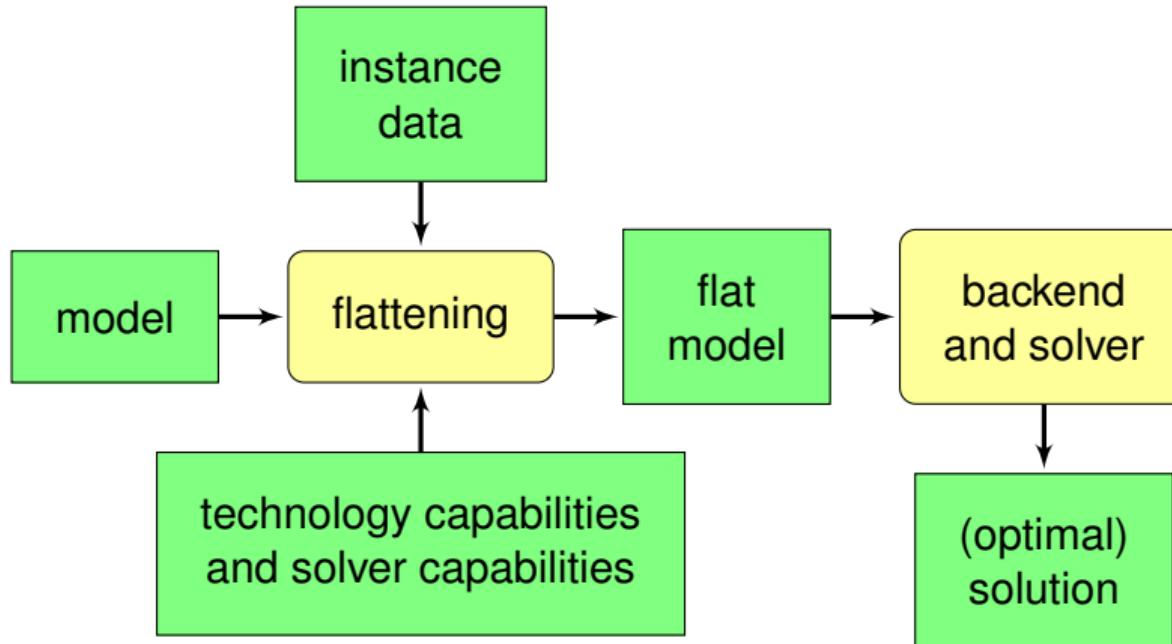


MiniZinc Challenge 2024: 15 (of 20) Problems and Winners

Problem and Model	Backend and Solver	Solving Technology
aircraft-disassembly	CP-SAT (by Google)	portfolio: LCG, MIP, CBLS
cable-tree-wiring	CP-SAT (by Google)	portfolio: LCG, MIP, CBLS
community-detection	MZN/Gurobi	MIP
compression	PicatSAT	SAT
concert-hall-cap	MZN/Gurobi	MIP
fox-geese-corn	MZN/Gurobi	MIP
hoist-benchmark	Chuffed	hybrid: LCG = CP + SAT
monitor-placement-1id	CP Optimizer (by IBM)	CP
neighbours	CP-SAT (by Google)	portfolio: LCG, MIP, CBLS
peaceable_queens	Gecode-Dexter	portfolio: CP, LNS
portal	Chuffed	hybrid: LCG = CP + SAT
tiny-cvrp	MZN/Gurobi	MIP
train-scheduling	CP-SAT (by Google)	portfolio: LCG, MIP, CBLS
word-equations	PicatSAT	SAT
yumi-dynamic	CP-SAT (by Google)	portfolio: LCG, MIP, CBLS



MiniZinc: Model Once, Solve Everywhere!



From a **single** language, one has access transparently to a wide range of solving technologies from which to choose.



There Is No Need to Reinvent the Wheel!

Before solving, each decision variable of a **type** that is non-native to the targeted solver is replaced by decision variables of native types, using some well-known linear / clausal / ... encoding.

Example (SAT)

The **order encoding** of integer decision variable `var 4..6: x` is

```
array[4..7] of var bool: B; % B[i] denotes truth of x >= i
constraint B[4];           % lower bound on x
constraint not B[7];       % upper bound on x
constraint B[4] \vee not B[5]; % consistency
constraint B[5] \vee not B[6]; % consistency
constraint B[6] \vee not B[7]; % consistency
```

For an integer decision variable with n domain values, there are $n + 1$ Boolean decision variables and n clauses, all 2-ary.



Before solving, each use of a non-native **predicate** or **function** is replaced by:

- either: its MiniZinc-provided default definition,
stated in terms of a kernel of imposed predicates;

Example (default; not to be used for IP and MIP)

`all_different ([x, y, z])` gives $x \neq y \wedge y \neq z \wedge z \neq x$.

- or: a backend-provided solver-specific definition,
using some well-known linear / clausal / ... encoding.

Example (IP and MIP)

A compact linearisation of $x \neq y$ is

```
var 0..1: p; % p = 1 denotes that x < y holds
int: Mx = ub(x-y+1); int: My = ub(y-x+1); % big-M constants
constraint x + 1 <= y + Mx * (1-p); % either x < y and p = 1
constraint y + 1 <= x + My * p; % or x > y and p = 0
```

One cannot naturally model graph colouring in IP,
but the problem has integer decision variables (ranging over the colours).



Benefits of Model-and-Solve with MiniZinc

- + Try many solvers of many technologies from 1 model.
- + A model improves with the state of the art of backends:
 - Type of decision variable: native representation or encoding.
 - Predicate: **reasoning**, **relaxation**, and definition.
 - Implementation of a solving technology.
- More on this in Topic 7: Solving Technologies.
- + For most managers, engineers, and scientists, it is easier with such a model-once-and-solve-everywhere toolchain to achieve good solution quality and high solving speed, including for harder data, and this without knowing (deeply) how the solvers work, compared to programming from first principles.



How to Solve a Constraint Problem?

1 Model the problem

Constraint
Problems

Combinatorial
Optimisation

Modelling
(in MiniZinc)

Solving

The MiniZinc
Toolchain

Course
Information

2 Solve the problem

Easy, right?



How to Solve a Constraint Problem?

1 Model the problem

- Understand the problem
- Choose the decision variables and their domains
- Choose predicates to formulate the constraints
- Formulate the objective function, if any
- Make sure the model really represents the problem
- Iterate!

2 Solve the problem

- Choose a solving technology
- Choose a backend
- Choose a search strategy, if not black-box search
- Improve the model
- Run the model and interpret the (lack of) solution(s)
- Debug the model, if need be
- Iterate!

Easy, right?



How to Solve a Constraint Problem?

1 Model the problem

- Understand the problem
- Choose the decision variables and their domains
- Choose predicates to formulate the constraints
- Formulate the objective function, if any
- Make sure the model really represents the problem
- Iterate!

2 Solve the problem

- Choose a solving technology
- Choose a backend
- Choose a search strategy, if not black-box search
- Improve the model
- Run the model and interpret the (lack of) solution(s)
- Debug the model, if need be
- Iterate!

Not so easy, but much easier than without a modelling tool!



Outline

1. Constraint Problems

2. Combinatorial Optimisation

3. Modelling (in MiniZinc)

4. Solving

5. The MiniZinc Toolchain

6. Course Information



Content

The use of tools for solving a combinatorial problem, by

- 1 first modelling it in a solving-technology-independent constraint-based modelling language, and
- 2 then running the model on an off-the-shelf solver.

We can now refine the course slogan:

How to solve combinatorial problems
without knowing how to solve combinatorial problems?



Learning Outcomes

In order to pass, the student must be able to:

- define the concept of combinatorial (optimisation or satisfaction) problem;
- explain the concept of constraint, as used in a constraint-based language;
- model a combinatorial problem in a solving-technology-independent constraint-based modelling language;
- compare empirically several models, say by introducing redundancy or by detecting and breaking symmetries;
- describe and compare solving technologies that can be used by the backends to a modelling language, including CP, LS, SAT, SMT, and MIP;
- choose suitable solving technologies for a new combinatorial problem, and motivate this choice;
- present and discuss topics related to the course content, orally and in writing, with a skill appropriate for the level of education.
 - ☞ written reports and oral resubmissions!



Organisation and *Suggested Time Budget*

Period 2: early November to mid January, budget = 133.3 h:

- No textbook: slides, [MiniZinc documentation](#), [Coursera](#) courses
- 1 **warm-up session** for learning the [MiniZinc](#) toolchain
- 3 teacher-chosen **assignments** with 3 **help sessions**, 1 **grading session**, and 1 **solution session** each, to be done in student-chosen duo team:
suggested budget = average of 21 hours / assignment / student (3 credits)
- 1 student-chosen **project**, with 8 **help sessions** (3 joint with Assignment 3), to be done in student-chosen duo team:
suggested budget = 49.5 hours / student (2 credits)
- 12 **lectures**, including a **mandatory** guest lecture: budget = 21 hours
- Prerequisites: basic concepts in algebra, combinatorics, logic, graph theory, set theory, and implementation of basic search algorithms



No Exam

Constraint
Problems

Combinatorial
Optimisation

Modelling
(in MiniZinc)

Solving

The MiniZinc
Toolchain

Course
Information

The course has no exam!

You must demonstrate — by writing **reports** — that you cannot only correctly and efficiently solve a constraint problem via a model,

but also motivate and explain your model in terms of **all** the course concepts, and experimentally demonstrate the correctness and efficiency of your model.



Lecture Topics

- Topic 1: Introduction
- Topic 2: Basic Modelling
- Topic 3: Constraint Predicates
- Topic 4: Modelling (for CP and LCG)
- Topic 5: Symmetry
- Topic 6: Case Studies
- Topic 7: Solving Technologies
- Topic 8: Reasoning & Search in CP & LCG
- (Topic 9: Modelling for CBLS)
- (Topic 10: Modelling for SAT, SMT, and OMT)
- (Topic 11: Modelling for MIP)



3 Assignment Cycles of 2 to 3 Weeks

Let D_i be the deadline day of Assignment i , with $i \in 1..3$:

- $D_i - 14$: publication and all needed material was taught: start!
- $D_i - 8$: help session a: participation strongly recommended!
- $D_i - 4$: help session b: participation strongly recommended!
- $D_i - 2$: help session c: participation strongly recommended!
- $D_i \pm 0$: submission, by 13:00 Swedish time on a Friday
- $D_i + 5$ by 16:00: initial score $a_i \in 0..5$ points
- $D_i + 6$: teamwise oral grading session for some $a_i \in \{1, 2\}$: possibility of earning 1 extra point for final score; otherwise final score = initial score
- $D_i + 6 = D_{i+1} - 8$: solution session and help session a



Assignments (3 credits) and Overall Grade

The final score on Assignment 1 is actually “pass” or “fail”.

Let $a_i \in 0..5$ be the **final score** on Assignment i , with $i \in 2..3$:

■ **20% threshold:** $\forall i \in 2..3 : a_i \geq 20\% \cdot 5 = 1$

No catastrophic failure on individual assignments

■ **50% threshold:** $a = a_2 + a_3 \geq 50\% \cdot (5 + 5) = 5$

The formula for the **assignment grade** in $3..5$ is at the course homepage

■ **Worth going full-blast:** An **assignment sum** $a \in 5..10$ is combined with a **project score** $p \in 5..10$ in order to determine the **overall grade** in $3..5$ according to a formula at the course homepage



Project (2 credits)

Topic:

- Model and solve a combinatorial problem that you are interested in, say for research, a course, a hobby, ...
- See the Project page at the course homepage for project ideas and the format for project proposals.

Deadlines, inevitably overlapping with Assignments 2 and 3:

- Wed 19 Nov 2025 at 13:00: upload several project proposals
- Wed 26 Nov 2025 at 13:00: secure our approval; start!
- Fri 19 Dec 2025 at 13:00: upload initial project report
- Fri 16 Jan 2026 at 13:00: upload final project report; get score $p \in 0..10$

50% threshold: $p \geq 50\% \cdot 10 = 5$

The formula for the **project grade** in 3..5 is at the course homepage



Project Guidelines

- Start early, despite the time overlap with Assignments 2 and 3.
- Attend the project help sessions, some jointly for Assignment 3.
- Read the Rules and Grading Criteria at the Project page.
- An **approach** is either a model for the entire problem, or a script (consider using **MiniZinc Python** and **JSON support**) with pre-processing + solving (possibly on a pipeline of multiple models) + post-processing: the final report is on **one** sufficiently complete and efficient approach.
- The initial report is on **one** approach, but it need be neither the final one, nor complete, nor efficient.



Project Guidelines (end)

- Model the constraints incrementally, and be prepared to backtrack to the choice of decision variables (aka viewpoint).
- If the instances are too easy, then you still need to demonstrate skills in the advanced concepts (49.5h!).
- If the instances are too hard, then relax the problem (say by some loss of precision on the objective value) or some instances (or both).
- Collaborate with other teams that work on the same problem for the parsing, generation, or simplification of shared instances, and so on (but *not* for modelling). There is *no* competition between such teams.
- Consider also using the powerful local-search backend Gecode-LNS for the experiments (see Assignment 3).



Assignment and Project Rules

Register a team by Sun 9 Nov 2025 at 23:59 at Studium:

- **Duo team:** Two consenting teammates sign up.
- **Solo team:** Apply in advance to the head teacher, who rarely agrees.
- **Random teammate?** Sign up at a Random Team; else you are bounced.

Other considerations:

- **Why (not) like this? Why no email reply?** See the FAQ list.
- **Teammate swapping:** Allowed, but to be declared to the helpdesk.
- **Teammate scores may differ** if no-show or passivity at grading session.
- **No freeloader:** Implicit honour declaration in reports that each teammate can individually explain everything; random checks will be made by us!
- **No plagiarism:** Implicit honour declaration in reports;
extremely powerful detection tools will be used by us;
suspected cases of using **or providing** must be reported!



How To Communicate by Email or Studium?

- If you have a question about the **lecture material or course organisation**, then email the head teacher. An immediate answer will be given right before and after lectures, as well as during their breaks.
- If you have a question about the **assignments or infrastructure**, then contact the assistants at a help session or solution session for an immediate answer.

Short *clarification* questions (that is: *not* about modelling or programming issues) that are either emailed (see the address at the course website) or posted (at the Studium discussion) to the **M4CO helpdesk** are answered as soon as possible during working days and hours.

No answer means that you should go to a help session:
almost all the assistants' budgeted time is allocated to grading and to the help, grading, and solution sessions.



What Has Changed Since Last Time?

Changes made by the TekNat Faculty:

- The course is moved to period 2, which is the requisite 10 weeks long (unlike the previous, shorter period 1).

Changes triggered by the formal and informal course evaluations:

- There are skeleton reports specific to the project and each assignment.
- To reduce the workload, peer reviewing an initial project report is dropped.
- Two additional project-only help sessions are scheduled, between the deadlines of the initial and final project reports.



What To Do Now?

- Bookmark and read the entire course website, especially its FAQ list.
- Read Sections 1 to 2.2 of the [MiniZinc Handbook](#).
- Get started on Assignment 1 and have questions ready for its first help session, which is already on Fri 7 Nov 2025.
- Register a duo team at Studium by Sun 9 Nov 2025 at 23:59, possibly upon advertising for a teammate at a course event or the discussion at Studium, and signing up at a Random Team at Studium as a last resort.
- Install the MiniZinc toolchain on your own hardware, if you have any.
- Be aware that few questions are tagged with MiniZinc at StackOverflow: you have to be prepared to read the documentation.